

```
import pandas as pd
```

```
df=pd.read_csv('dataset/cross-validation.csv')
```

```
df.head()
```

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	\
0	LP001002	Male	No	0	Graduate	No	
1	LP001003	Male	Yes	1	Graduate	No	
2	LP001005	Male	Yes	0	Graduate	Yes	
3	LP001006	Male	Yes	0	Not Graduate	No	
4	LP001008	Male	No	0	Graduate	No	

	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	\
0	5849	0.0	NaN	360.0	
1	4583	1508.0	128.0	360.0	
2	3000	0.0	66.0	360.0	
3	2583	2358.0	120.0	360.0	
4	6000	0.0	141.0	360.0	

	Credit_History	Property_Area	Loan_Status
0	1.0	Urban	Y
1	1.0	Rural	N
2	1.0	Urban	Y
3	1.0	Urban	Y
4	1.0	Urban	Y

```
df.describe()
```

	ApplicantIncome	CoapplicantIncome	LoanAmount
Loan_Amount_Term \			
count	614.000000	614.000000	592.000000
600.000000			
mean	5403.459283	1621.245798	146.412162
342.000000			
std	6109.041673	2926.248369	85.587325
65.12041			
min	150.000000	0.000000	9.000000
12.000000			
25%	2877.500000	0.000000	100.000000
360.000000			
50%	3812.500000	1188.500000	128.000000
360.000000			
75%	5795.000000	2297.250000	168.000000
360.000000			
max	81000.000000	41667.000000	700.000000
480.000000			

	Credit_History
count	564.000000
mean	0.842199

```
std      0.364878
min      0.000000
25%      1.000000
50%      1.000000
75%      1.000000
max      1.000000
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 614 entries, 0 to 613
```

```
Data columns (total 13 columns):
```

#	Column	Non-Null Count	Dtype
0	Loan_ID	614 non-null	object
1	Gender	601 non-null	object
2	Married	611 non-null	object
3	Dependents	599 non-null	object
4	Education	614 non-null	object
5	Self_Employed	582 non-null	object
6	ApplicantIncome	614 non-null	int64
7	CoapplicantIncome	614 non-null	float64
8	LoanAmount	592 non-null	float64
9	Loan_Amount_Term	600 non-null	float64
10	Credit_History	564 non-null	float64
11	Property_Area	614 non-null	object
12	Loan_Status	614 non-null	object

```
dtypes: float64(4), int64(1), object(8)
```

```
memory usage: 62.5+ KB
```

```
df.isnull().sum()
```

Loan_ID	0
Gender	13
Married	3
Dependents	15
Education	0
Self_Employed	32
ApplicantIncome	0
CoapplicantIncome	0
LoanAmount	22
Loan_Amount_Term	14
Credit_History	50
Property_Area	0
Loan_Status	0

```
dtype: int64
```

```
df.shape
```


```
(614, 13)
```

```
df['Gender'].fillna('Other',inplace=True)
df['Dependents'].fillna(0,inplace=True)
df['Self_Employed'].fillna('Both',inplace=True)
df['Credit_History'].fillna(0,inplace=True)

import seaborn as sns
import matplotlib.pyplot as plt


sns.countplot(x='Gender',data=df)

<Axes: xlabel='Gender', ylabel='count'>
```




```
sns.countplot(x='Education',hue='Gender',data=df)


<Axes: xlabel='Education', ylabel='count'>
```




```
s=df['Property_Area'].value_counts()/  
df['Property_Area'].value_counts().sum()*100  
plt.pie(s.values,labels=s.index,autopct='%1.1f%%')  
plt.show()
```



```
sns.lineplot(x='ApplicantIncome',y='LoanAmount',data=df)  
<Axes: xlabel='ApplicantIncome', ylabel='LoanAmount'>
```



```
sns.lineplot(x='CoapplicantIncome',y='LoanAmount',data=df)  
<Axes: xlabel='CoapplicantIncome', ylabel='LoanAmount'>
```



```
import numpy as np

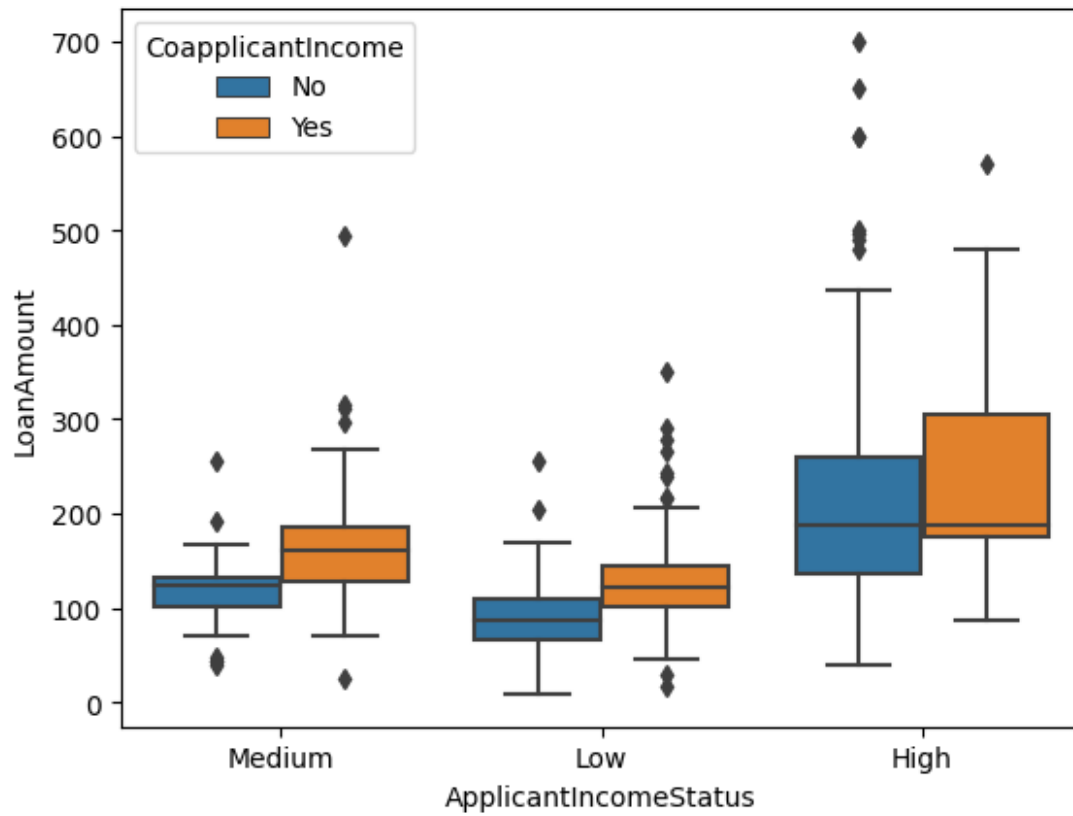
incomeType=[]
for i in df['ApplicantIncome']:
    if i>6000:
        incomeType.append('High')
    elif i<6000 and i>4000:
        incomeType.append('Medium')
    else:
        incomeType.append('Low')
df['ApplicantIncomeStatus']=pd.Series(incomeType)
df['CoapplicantIncome']=np.where(df['CoapplicantIncome']>0,'Yes','No')

df['ApplicantIncomeStatus'].value_counts()


ApplicantIncomeStatus
Low      340
Medium   137
High     137
Name: count, dtype: int64

sns.boxplot(x='ApplicantIncomeStatus',y='LoanAmount',hue='CoapplicantIncome',data=df)

<Axes: xlabel='ApplicantIncomeStatus', ylabel='LoanAmount'>
```



```
sns.boxplot(x='Property_Area',y='LoanAmount',data=df)  
<Axes: xlabel='Property_Area', ylabel='LoanAmount'>
```

```
df.groupby(['Married', 'Education', 'Self_Employed', 'Dependents'])
['LoanAmount'].mean()
```

Married	Education	Self_Employed	Dependents	
No	Graduate	Both	0	174.833333
			1	104.000000
			2	120.000000
			3+	350.000000
	Not Graduate	Both	0	123.000000
			1	132.000000
	Graduate	No	0	94.137931
			1	91.666667
			3+	148.000000
Yes	Graduate	Both	0	173.000000

				1	170.666667
				2	102.000000
				3+	110.000000
				0	172.500000
				0	147.948276
				1	167.468085
				2	153.058824
				3+	216.074074
				0	70.000000
				0	149.916667
				1	203.692308
				2	223.307692
				3+	182.333333
				0	125.000000
				2	98.000000
				0	126.666667
				0	128.807692
				1	146.692308
				2	115.300000
				3+	103.111111
				0	138.000000
				0	121.500000
				2	142.000000
				3+	128.333333

Name: LoanAmount, dtype: float64

df

	Loan_ID	Gender	Married	Dependents	Education	
Self_Employed	\					
0	LP001002	Male	No	0	Graduate	No
1	LP001003	Male	Yes	1	Graduate	No
2	LP001005	Male	Yes	0	Graduate	Yes
3	LP001006	Male	Yes	0	Not Graduate	No
4	LP001008	Male	No	0	Graduate	No
..
609	LP002978	Female	No	0	Graduate	No
610	LP002979	Male	Yes	3+	Graduate	No
611	LP002983	Male	Yes	1	Graduate	No
612	LP002984	Male	Yes	2	Graduate	No
613	LP002990	Female	No	0	Graduate	Yes

	ApplicantIncome	CoapplicantIncome	LoanAmount	
Loan_Amount_Term \				
0	5849	No	NaN	360.0
1	4583	Yes	128.0	360.0
2	3000	No	66.0	360.0
3	2583	Yes	120.0	360.0
4	6000	No	141.0	360.0
..
609	2900	No	71.0	360.0
610	4106	No	40.0	180.0
611	8072	Yes	253.0	360.0
612	7583	No	187.0	360.0
613	4583	No	133.0	360.0

	Credit_History	Property_Area	Loan_Status	ApplicantIncomeStatus
0	1.0	Urban	Y	Medium
1	1.0	Rural	N	Medium
2	1.0	Urban	Y	Low
3	1.0	Urban	Y	Low
4	1.0	Urban	Y	Low
..
609	1.0	Rural	Y	Low
610	1.0	Rural	Y	Medium
611	1.0	Urban	Y	High
612	1.0	Urban	Y	High
613	0.0	Semiurban	N	Medium

[614 rows x 14 columns]

```
df['Gender'] = df['Gender'].map({'Male': 0, 'Female': 1, 'Other': 2})
df['Married'] = df['Married'].map({'Yes': 0, 'No': 1})
df['Dependents'] = df['Dependents'].map({'0': 0, '1': 1, '2': 2, '3+': 3})
df['Education'] = df['Education'].map({'Graduate': 0, 'Not Graduate': 1})
df['Self_Employed'] = df['Self_Employed'].map({'Yes': 0, 'No': 1, 'Both': 2})
```

```

df['Property_Area'] = df['Property_Area'].map({'Urban': 0, 'Rural': 1, 'Semiurban': 2})

df['Loan_Status'] = df['Loan_Status'].map({'Y': 0, 'N': 1})
df['ApplicantIncomeStatus'] =
df['ApplicantIncomeStatus'].map({'Low':0, 'Medium':1, 'High':2})

df['CoapplicantIncome']=df['CoapplicantIncome'].map({'Yes':0, 'No':1})

df['Loan_Amount_Term'] =
df['Loan_Amount_Term'].replace(np.nan,df['Loan_Amount_Term'].mean())
df['LoanAmount'] =
df['LoanAmount'].replace(np.nan,df['LoanAmount'].mean())
df['Dependents'] = df['Dependents'].replace(np.nan,0)
df['Married'] = df['Married'].replace(np.nan,1)

df.isna().sum()

Loan_ID          0
Gender           0
Married          0
Dependents       0
Education        0
Self_Employed   0
ApplicantIncome  0
CoapplicantIncome 0
LoanAmount       0
Loan_Amount_Term 0
Credit_History  0
Property_Area    0
Loan_Status      0
ApplicantIncomeStatus 0
dtype: int64

columns = df.select_dtypes(exclude='int').columns
columns

Index(['Loan_ID', 'Married', 'Dependents', 'LoanAmount',
      'Loan_Amount_Term',
      'Credit_History'],
      dtype='object')

df=df.drop(['Loan_ID'],axis=1)
df=df.drop(['ApplicantIncome'],axis=1)

df['Loan_Amount_Term']=df['Loan_Amount_Term'].div(max(df['Loan_Amount_Term']))

df['LoanAmount']=df['LoanAmount'].div(max(df['LoanAmount']))

df

```

	Gender	Married	Dependents	Education	Self_Employed
CoapplicantIncome \					
0	0	1.0	0.0	0	1
1					
1	0	0.0	1.0	0	1
0					
2	0	0.0	0.0	0	0
1					
3	0	0.0	0.0	1	1
0					
4	0	1.0	0.0	0	1
1					
...
...					
609	1	1.0	0.0	0	1
1					
610	0	0.0	3.0	0	1
1					
611	0	0.0	1.0	0	1
0					
612	0	0.0	2.0	0	1
1					
613	1	1.0	0.0	0	0
1					
	LoanAmount	Loan_Amount_Term	Credit_History	Property_Area	
Loan_Status \					
0	0.209160	0.750	1.0	0	
0					
1	0.182857	0.750	1.0	1	
1					
2	0.094286	0.750	1.0	0	
0					
3	0.171429	0.750	1.0	0	
0					
4	0.201429	0.750	1.0	0	
0					
...	
...					
609	0.101429	0.750	1.0	1	
0					
610	0.057143	0.375	1.0	1	
0					
611	0.361429	0.750	1.0	0	
0					
612	0.267143	0.750	1.0	0	
0					
613	0.190000	0.750	0.0	2	
1					

```

ApplicantIncomeStatus
0      1
1      1
2      0
3      0
4      0
..    ..
609    0
610    1
611    2
612    2
613    1

```


[614 rows x 12 columns]

```

plt.figure(figsize=[15,8])
sns.heatmap(df.corr(),annot=True)

```

<Axes: >



```

X = df.drop('Loan_Status', axis=1)
Y = df['Loan_Status']

```

```

from sklearn.model_selection import train_test_split
x_train , x_test , y_train , y_test = train_test_split(X, Y ,
test_size = 0.2 , random_state = 42, stratify=Y)

from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
x_train_scaled = scaler.fit_transform(x_train)
x_test_scaled = scaler.transform(x_test)

from sklearn.linear_model import LogisticRegression
model = LogisticRegression()
model.fit(x_train_scaled,y_train)
y_pred = model.predict(x_test_scaled)

from sklearn.metrics import classification_report
print(classification_report(y_test,y_pred))

```

	precision	recall	f1-score	support
0	0.80	0.85	0.82	85
1	0.61	0.53	0.56	38
accuracy			0.75	123
macro avg	0.70	0.69	0.69	123
weighted avg	0.74	0.75	0.74	123

K-Fold Validation

```

X=df.drop('Loan_Status', axis=1)
Y= df['Loan_Status']
nData = int(X.shape[0]/5)
folds=[]
targets=[]
s=0
e=nData
while(e<X.shape[0]):
    fold=X.loc[s:e]
    target=Y.loc[s:e]
    folds.append(fold)
    targets.append(target)
    s=e+1
    e=s+nData
folds.append(X.loc[s:])
targets.append(Y.loc[s:])

train_data=[]
target_train=[]
test_data=[]

```

```

target_test=[]
k=len(folds)
for i in range(0,k):
    test_data.append(folds[i])
    target_test.append(targets[i])
    data=[]
    target=[]
    for j in range(0,k):
        if(i==j):
            continue
        data.append(folds[j])
        target.append(targets[j])
    train_data.append(pd.concat(data))
    target_train.append(pd.concat(target))

modell = LogisticRegression(penalty=None, solver='saga', max_iter=300)

for i in range(0,k):
    modell.fit(train_data[i],target_train[i])
    y_pred = modell.predict(test_data[i])
    print(f"Trained with Fold {i+1}")
    print(classification_report(target_test[i],y_pred))
    print(f"Accuracy: {modell.score(test_data[i], target_test[i])}")
    print("-----")

```

Trained with Fold 1

	precision	recall	f1-score	support
0	0.77	0.93	0.84	83
1	0.74	0.42	0.54	40
accuracy			0.76	123
macro avg	0.75	0.68	0.69	123
weighted avg	0.76	0.76	0.74	123

Accuracy: 0.7642276422764228

Trained with Fold 2

	precision	recall	f1-score	support
0	0.76	0.88	0.81	84
1	0.60	0.38	0.47	39
accuracy			0.72	123
macro avg	0.68	0.63	0.64	123
weighted avg	0.71	0.72	0.70	123

Accuracy: 0.7235772357723578

Trained with Fold 3

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

0	0.78	0.87	0.83	87
1	0.58	0.42	0.48	36

accuracy			0.74	123
macro avg	0.68	0.65	0.65	123
weighted avg	0.72	0.74	0.73	123

Accuracy: 0.7398373983739838

Trained with Fold 4

	precision	recall	f1-score	support
0	0.82	0.87	0.85	85
1	0.67	0.58	0.62	38

accuracy			0.78	123
macro avg	0.74	0.72	0.73	123
weighted avg	0.77	0.78	0.78	123

Accuracy: 0.7804878048780488

Trained with Fold 5

	precision	recall	f1-score	support
0	0.80	0.93	0.86	83
1	0.77	0.51	0.62	39

accuracy			0.80	122
macro avg	0.79	0.72	0.74	122
weighted avg	0.79	0.80	0.78	122

Accuracy: 0.7950819672131147
